

```
from transformers import pipeline, set_seed, GPT2Tokenizer
import os
import nltk
```

```
WARNING:torchao.kernel.intmm:Warning: Detected no triton, on systems without Triton certain kernels will not work
```

```
file_path = "unit 1.txt"
```

```
try:
    with open(file_path,'r',encoding="utf-8") as f:
        text=f.read()
    print("File loaded successfully!")
except FileNotFoundError:
    print(f"Error: {file_path} not found!")
```

```
File loaded successfully!
```

```
print("--- Data Preview ---")
print(text[:500]+"...")
```

```
--- Data Preview ---
Generative AI and Its Applications: A Foundational Briefing
```

```
Executive Summary
```

```
This document provides a comprehensive overview of Generative AI, synthesizing foundational concepts, technological underpinnings, and practical applications.
```

```
set_seed(42);
```

```
prompt = "Generative AI is an elective subject"
```

```
fast_generator = pipeline("text-generation", model="distilgpt2")
output_fast = fast_generator(prompt, max_length=50, num_return_sequences=1)
print(output_fast[0]['generated_text'])
```

```
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as an environment variable, or pass it directly to the API.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
    warnings.warn(
config.json: 100%                                         762/762 [00:00<00:00, 53.9kB/s]
model.safetensors: 100%                                     353M/353M [00:07<00:00, 32.4MB/s]
generation_config.json: 100%                                124/124 [00:00<00:00, 1.49kB/s]
tokenizer_config.json: 100%                                 26.0/26.0 [00:00<00:00, 320B/s]
vocab.json: 100%                                         1.04M/1.04M [00:00<00:00, 5.13MB/s]
merges.txt: 100%                                         456k/456k [00:00<00:00, 10.7MB/s]
tokenizer.json: 100%                                     1.36M/1.36M [00:00<00:00, 18.3MB/s]

Device set to use cpu
Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly activate truncation.
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
Both `max_new_tokens` (=256) and `max_length` (=50) seem to have been set. `max_new_tokens` will take precedence. Please refer to the documentation for more information.
Generative AI is an elective subject that is currently a major focus in AI research.
```

```
fast_generator = pipeline("text-generation", model="gpt2")
output_fast = fast_generator(prompt, max_length=50, num_return_sequences=1)
print(output_fast[0]['generated_text'])
```

```
config.json: 100% 665/665 [00:00<00:00, 13.8kB/s]
model.safetensors: 100% 548M/548M [00:08<00:00, 46.0MB/s]
generation_config.json: 100% 124/124 [00:00<00:00, 2.69kB/s]
tokenizer_config.json: 100% 26.0/26.0 [00:00<00:00, 561B/s]
vocab.json: 100% 1.04M/1.04M [00:00<00:00, 12.1MB/s]
merges.txt: 100% 456k/456k [00:00<00:00, 9.67MB/s]
tokenizer.json: 100% 1.36M/1.36M [00:00<00:00, 24.2MB/s]
Device set to use cpu
Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly activate truncation
tokenizer = GPT2Tokenizer.from_pretrained("gpt2")
Generative AI is an elective subject, i.e. an AI that can learn from a specific situation. Once you have created an agent, you can use it to generate text.
sample_sentence = "Transformers have encoders and decoders"
```

```
tokens = tokenizer.tokenize(sample_sentence)
print(f"Tokens: {tokens}")
```

```
Tokens: ['Transformers', 'have', 'encoders', 'and', 'decoders']
```

```
token_ids = tokenizer.convert_tokens_to_ids(tokens)
print(f"Token IDs: {token_ids}")
```

```
Token IDs: [41762, 364, 423, 2207, 375, 364, 290, 875, 375, 364]
```

```
nltk.download('averaged_perceptron_tagger_eng', quiet=True)
nltk.download('punkt', quiet=True)
nltk.download('punkt_tab')
```

```
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data]  Unzipping tokenizers/punkt_tab.zip.
True
```

```
pos_tags = nltk.pos_tag(nltk.word_tokenize(sample_sentence))
print(f"POS Tags: {pos_tags}")
```

```
POS Tags: [('Transformers', 'NNS'), ('have', 'VBP'), ('encoders', 'NNS'), ('and', 'CC'), ('decoders', 'NNS')]
```

```
ner_pipeline = pipeline("ner", model="dbmdz/bert-large-cased-finetuned-conll03-english")
```

```
config.json: 100% 998/998 [00:00<00:00, 55.4kB/s]
model.safetensors: 100% 1.33G/1.33G [00:16<00:00, 75.5MB/s]
Some weights of the model checkpoint at dbmdz/bert-large-cased-finetuned-conll03-english were not used when initializing BertForTokenClassification - This IS expected if you are initializing BertForTokenClassification from the checkpoint of a model trained on another task - This IS NOT expected if you are initializing BertForTokenClassification from the checkpoint of a model that you expect to tokenizer_config.json: 100% 60.0/60.0 [00:00<00:00, 3.50kB/s]
vocab.txt: 213k? [00:00<00:00, 11.8MB/s]
Device set to use cpu
```

```
snippet = text[:1000]
entities = ner_pipeline(snippet)

print(f"{'Entity':<20} | {'Type':<10} | {'Score':<5}")
print("-" * 45)

for entity in entities:
    if entity['score'] > 0.9:
        print(f"{entity['word']:<20} | {entity['entity']:<10} | {entity['score']:.2f}")
```

Entity	Type	Score
AI	I-MISC	0.98
AI	I-MISC	0.99
P	I-ORG	1.00
##ES	I-ORG	1.00
University	I-ORG	0.97
AI	I-MISC	0.98
Language	I-MISC	0.98
Models	I-MISC	0.99
LL	I-MISC	0.95
GP	I-MISC	0.96
#T	I-MISC	0.93

4	I-MISC	0.97
Trans	I-MISC	0.99
##former	I-MISC	0.99

```
transformer_section = """
The introduction of the Transformer architecture in the 2017 paper "Attention is all you need" was a watershed moment in AI
The fundamental innovation of the Transformer is the attention mechanism. This component allows the model to weigh the importance of different words in a sentence based on their context.
The Transformer architecture consists of an encoder stack (to process the input) and a decoder stack (to generate the output).
"""


```

```
fast_sum = pipeline("summarization", model="sshleifer/distilbart-cnn-12-6")
res_fast = fast_sum(transformer_section, max_length=60, min_length=30, do_sample=False)
print(res_fast[0]['summary_text'])

Device set to use cpu
The introduction of the Transformer architecture in the 2017 paper "Attention is all you need" was a watershed moment in AI
```

```
smart_sum = pipeline("summarization", model="facebook/bart-large-cnn")
res_smart = smart_sum(transformer_section, max_length=60, min_length=30, do_sample=False)
print(res_smart[0]['summary_text'])

1 you need" was a watershed moment in AI. It provided a more effective and scalable way to handle sequential data like text.
```

```
qa_pipeline = pipeline("question-answering", model="distilbert-base-cased-distilled-squad")
```

```
config.json: 100% 473/473 [00:00<00:00, 33.8kB/s]
model.safetensors: 100% 261M/261M [00:06<00:00, 39.2MB/s]
tokenizer_config.json: 100% 49.0/49.0 [00:00<00:00, 2.03kB/s]
vocab.txt: 100% 213k/213k [00:00<00:00, 9.73MB/s]
tokenizer.json: 100% 436k/436k [00:00<00:00, 21.4MB/s]
Device set to use cpu
```

```
questions = [
    "What is the fundamental innovation of the Transformer?",
    "What are the risks of using Generative AI?"
]

for q in questions:
    res = qa_pipeline(question=q, context=text[:5000])
    print(f"\nQ: {q}")
    print(f"A: {res['answer']}")
```

Q: What is the fundamental innovation of the Transformer?
A: to identify hidden patterns, structures, and relationships within the data

Q: What are the risks of using Generative AI?
A: data privacy, intellectual property, and academic integrity

```
mask_filler = pipeline("fill-mask", model="bert-base-uncased")
```

```
config.json: 100% 570/570 [00:00<00:00, 43.9kB/s]
model.safetensors: 100% 440M/440M [00:18<00:00, 21.8MB/s]
Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForMaskedLM: ['bert.pooler.dense']
- This IS expected if you are initializing BertForMaskedLM from the checkpoint of a model trained on another task or with a different architecture.
- This IS NOT expected if you are initializing BertForMaskedLM from the checkpoint of a model that you expect to be exactly identical.
tokenizer_config.json: 100% 48.0/48.0 [00:00<00:00, 3.15kB/s]
vocab.txt: 100% 232k/232k [00:00<00:00, 11.3MB/s]
tokenizer.json: 100% 466k/466k [00:00<00:00, 24.8MB/s]
Device set to use cpu
```

```
masked_sentence = "The goal of Generative AI is to create new [MASK]."
preds = mask_filler(masked_sentence)

for p in preds:
    print(f"{p['token_str']}: {p['score']:.2f}")

applications: 0.06
ideas: 0.05
problems: 0.05
```

```
systems: 0.04
information: 0.03
```